

Metadata

Course: DS 5001
Module: 02 Homework KEY
Topic: Text Models
Author: R.C. Alvaraddo
Date: 30 January 2023 (revised)

Overview

Students parse a second text following the pattern of the first, and then combine them to create a corpus.

With the corpus, students observe basic descriptive statistical features.

Instructions

In this exercise, you will convert a different text from raw text into a data frame of tokens and preserving its OHCO. Then you will extract some statistical features from the resulting corpus.

Follow these instructions:

1. Download the attached Gutenberg version of Jane Austen's *Sense and Sensibility* (`pg161.txt`).
2. Create a notebook to convert the raw text into a data frame of tokens, just as we did with *Persuasion*. You may use the notebook from the lab as your guide.
3. Specifically, make sure you complete these tasks:
 - (a) Remove Gutenberg's front and back matter using the lines that indicate the start and end of the project.
 - (b) Chunk by chapter, using the pattern of locating the headers in the data frame, assigning them numbers, forward-filling those numbers, and then grouping by number (and cleaning up).
 - (c) Split resulting data frame into paragraphs using the regex provided.
 - (d) Split resulting data frame into sentences using the regex provided.
 - (e) Split resulting data frame into tokens using the regex provided.
 - (f) Be sure to include the OHCO of Chapters, Paragraphs, and Sentences in your data frame's index.
4. Once you have done this, combine both *Persuasion* and *Sense and Sensibility* into a single data frame with an appropriately modified OHCO list. In other words, make sure your index includes a new index level for the book. Use the attached CSV (`austen-persuasion.csv`) to get the *Persuasion* data and then import it into your notebook as a data frame.
5. From the combined data frame, extract a vocabulary, i.e. a data frame with term string as index, along with term frequency and term length as features.

6. After you have done all this, answer the following questions by extracting features from the corpus.
- (a) How many raw tokens are in the combined data frame?
 - (b) How many distinct terms are there in the combined data frame (i.e. how big is the vocabulary)?
 - (c) How many more terms does the vocabulary of *Sense and Sensibility* have than that of *Persuasion*?
 - (d) What is the average number of tokens, rounded to an integer, per chapter in the corpus?
 - (e) What is the average number of tokens, rounded to an integer, per paragraph in the corpus?

Summary

- Convert `pg161.txt` into `austen-sense.csv` with OHCO of chapters, paragraphs, sentences, and tokens.
- Combine tokenized dataframes of `austen-sense.csv` and `austen-persuasion.csv` into `austen-combo.csv`.
- Extract a vocabulary with term frequencies.
- Answer the questions.

Set Up

```
import pandas as pd
import seaborn as sns

sns.set()

data_home = '../.../repo/lessons/data'

text_file = f"{data_home}/gutenberg/pg161.txt"
csv_file1 = f"{data_home}/output/austen-sense.csv" # To be created
csv_file2 = f"{data_home}/output/austen-persuasion.csv" # Already created
csv_combo = f"{data_home}/output/austen-combo.csv" # To be created
OHCO = ['chap_num', 'para_num', 'sent_num', 'token_num']
```

Import Text

Import *Sense and Sensibility* into a dataframe.

```
LINES = pd.DataFrame(open(text_file, 'r', encoding='utf-8-sig').readlines(),
                      columns=['line_str'])
LINES.index.name = 'line_num'
LINES.line_str = LINES.line_str.str.strip()

LINES.sample(10)
```

```

line_str
line_num
3690         they say it is a sweet pretty place."
9395     since it can advance him so little towards wha...
5744     arrival, without once stirring from her seat, ...
12762     1.E.1. The following sentence, with active li...
8713     barbarous have I been to you!--you, who have b...
5656         of her real situation with respect to him.
5911     circumstances, it was better for both that the...
2804     Mrs. Dashwood was sorry for what she had said;...
3870     "Is Mr. Willoughby much known in your part of ...
2501     will always be welcome; for I will not press y...

```

Extract title of work from first line

```

title = LINES.loc[0].line_str.replace('The Project Gutenberg EBook of ', '')
title
'Sense and Sensibility, by Jane Austen'
LINES.head()

```

```

line_str
line_num
0         The Project Gutenberg EBook of Sense and Sensi...
1
2         This eBook is for the use of anyone anywhere a...
3         almost no restrictions whatsoever. You may co...
4         re-use it under the terms of the Project Guten...

```

Remove Gutenberg's front and back matter

```

a = LINES.line_str.str.match(r"\*.*\s*START OF (THE|THIS) PROJECT")
b = LINES.line_str.str.match(r"\*.*\s*END OF (THE|THIS) PROJECT")
an = LINES.loc[a].index[0]
bn = LINES.loc[b].index[0]
LINES = LINES.loc[an + 1 : bn - 2]
LINES

```

```

line_str
line_num
20
21
22

```

```

23
24
...
12661
12662
12663
12664
12665      End of the Project Gutenberg EBook of Sense an...

[12646 rows x 1 columns]

```

Chunk by chapter

Find all chapter headers

```

chap_lines = LINES.line_str.str.match(r"^\s*(chapter|letter)\s+(\d+)", case=False)
LINES.loc[chap_lines]

      line_str
line_num
42      CHAPTER 1
196     CHAPTER 2
399     CHAPTER 3
561     CHAPTER 4
756     CHAPTER 5
858     CHAPTER 6
986     CHAPTER 7
1112    CHAPTER 8
1244    CHAPTER 9
1448    CHAPTER 10
1665    CHAPTER 11
1816    CHAPTER 12
1997    CHAPTER 13
2281    CHAPTER 14
2440    CHAPTER 15
2718    CHAPTER 16
2945    CHAPTER 17
3153    CHAPTER 18
3331    CHAPTER 19
3632    CHAPTER 20
3913    CHAPTER 21
4214    CHAPTER 22
4532    CHAPTER 23
4767    CHAPTER 24
5001    CHAPTER 25

```

```

5197      CHAPTER 26
5454      CHAPTER 27
5732      CHAPTER 28
5883      CHAPTER 29
6324      CHAPTER 30
6628      CHAPTER 31
7004      CHAPTER 32
7278      CHAPTER 33
7601      CHAPTER 34
7888      CHAPTER 35
8152      CHAPTER 36
8456      CHAPTER 37
8900      CHAPTER 38
9205      CHAPTER 39
9408      CHAPTER 40
9706      CHAPTER 41
9977      CHAPTER 42
10155     CHAPTER 43
10490     CHAPTER 44
11060     CHAPTER 45
11278     CHAPTER 46
11571     CHAPTER 47
11838     CHAPTER 48
11986     CHAPTER 49
12410     CHAPTER 50

```

Assign numbers to chapters

```

chap_nums = [i+1 for i in range(LINES.loc[chap_lines].shape[0])]
LINES.loc[chap_lines, 'chap_num'] = chap_nums

```

Forward-fill chapter numbers to following text lines

```

LINES.chap_num = LINES.chap_num.ffill()

```

Clean up

```

# LINES = LINES.loc[~LINES.chap_num.isna()] # Remove chapter heading lines
LINES = LINES.dropna(subset=['chap_num'])
LINES = LINES.loc[~chap_lines] # Remove everything before Chapter 1
LINES.chap_num = LINES.chap_num.astype('int') # Convert chap_num from float to int
LINES.sample(10)

line_str  chap_num

line_num

```

3994	demands which this politeness made on it, was ...	21
10943	happy, and afterwards returned to town to be g...	44
8132		35
7473		33
1436	destroyed all its ingenuity."	9
602	perceive how you could express yourself more w...	4
7157	Mrs. Jennings, who knew nothing of all this, w...	32
11596	In the evening, when they were all three toget...	47
10595		44
12002		49

Group lines by chapter num

```
CHAPS = LINES.groupby(OHCO[:1]).line_str.apply(lambda x: '\n'.join(x)).to_frame('chap_str')
CHAPS.head()
```

	chap_str
chap_num	
1	\n\nThe family of Dashwood had long been settl...
2	\n\nMrs. John Dashwood now installed herself m...
3	\n\nMrs. Dashwood remained at Norland several ...
4	\n\n"What a pity it is, Elinor," said Marianne...
5	\n\nNo sooner was her answer dispatched, than ...

Split into paragraphs

```
PARAS = CHAPS['chap_str'].str.split(r'\n\n+', expand=True).stack()\
.to_frame('para_str')
PARAS.index.names = OHCO[:2]
PARAS.head()
```

		para_str
chap_num	para_num	
1	0	
	1	The family of Dashwood had long been settled i...
	2	By a former marriage, Mr. Henry Dashwood had o...
	3	The old gentleman died: his will was read, and...
	4	Mr. Dashwood's disappointment was, at first, s...

```
PARAS['para_str'] = PARAS['para_str'].str.replace(r'\n', ' ', regex=True).str.strip()
PARAS = PARAS[~PARAS['para_str'].str.match(r'^\s*$')] # Remove empty paragraphs
PARAS.head()
```

		para_str
chap_num	para_num	
1	1	The family of Dashwood had long been settled i...

```

2      By a former marriage, Mr. Henry Dashwood had o...
3      The old gentleman died: his will was read, and...
4      Mr. Dashwood's disappointment was, at first, s...
5      His son was sent for as soon as his danger was...

```

Split into sentences

NOTE: ADDED " to regex in split()

```

SENTS = PARAS['para_str'].str.split(r'[\.?!;:"]+', expand=True).stack()\
    .to_frame().rename(columns={0:'sent_str'})
SENTS.index.names = OHCO[:3]
SENTS = SENTS[~SENTS['sent_str'].str.match(r'^\s*$')] # Remove empty paragraphs
SENTS.sent_str = SENTS.sent_str.str.strip()

SENTS.head()

```

chap_num	para_num	sent_num	sent_str
1	1	0	The family of Dashwood had long been settled i...
		1	Their estate was large, and their residence wa...
		2	The late owner of this estate was a single man...
		3	But her death, which happened ten years before...
		4	for to supply her loss, he invited and receive...

Split into tokens

```

TOKENS = SENTS['sent_str'].str.split(r"[\s',-]+", expand=True).stack()\
    .to_frame('token_str')
TOKENS.index.names = OHCO[:4]

TOKENS['term_str'] = TOKENS.token_str.str.replace(r"[\W_]+", '', regex=True).str.lower()

TOKENS

```

chap_num	para_num	sent_num	token_num	token_str	term_str
1	1	0	0	The	the
			1	family	family
			2	of	of
			3	Dashwood	dashwood
			4	had	had
...		
50	23	0	8	and	and
			9	Sensibility	sensibility
			10	by	by
			11	Jane	jane

```
[122257 rows x 2 columns]
```

Save work to CSV

```
TOKENS.to_csv(csv_file1)
```

Combine the two into a Corpus

```
csv_file2 = f"{data_home}/output/austen-persuasion.csv"
df1 = pd.read_csv(csv_file1)
df2 = pd.read_csv(csv_file2)

len(df1), len(df2)
(122257, 85014)

df1['book_id'] = 1 # They may use the string for the titles here
df2['book_id'] = 2

LIB = {
    1: 'Sense & Sensibility',
    2: 'Persuasion'
}

CORPUS = pd.concat([df1, df2])
OHC02 = ['book_id'] + OHCO
CORPUS = CORPUS.set_index(OHC02)
# CORPUS.sample(10)

len(CORPUS), CORPUS.shape[0], CORPUS.token_str.count()
(207271, 207271, 205599)
```

Extract a vocabulary V

```
CORPUS['term_str'] = CORPUS.token_str.str.replace(r"\W+", "", regex=True).str.lower()
V = CORPUS.term_str.value_counts().to_frame('n')
V.index.name = 'term_str'
V['n_chars'] = V.index.str.len()

len(V)
8239

V.n_chars.mean()
```


7.5543148440344705

CORPUS

book_id	chap_num	para_num	sent_num	token_num	token_str	term_str
1	1	1	0	0	The	the
				1	family	family
				2	of	of
				3	Dashwood	dashwood
				4	had	had
...				
2	24	13	0	6	of	of
				7	Persuasion	persuasion
				8	by	by
				9	Jane	jane
				10	Austen	austen

[207271 rows x 2 columns]

V

	n	n_chars
term_str		
the	7435	3
to	6923	2
and	6290	3
of	6146	2
her	3747	3
...
unconquerable	1	13
outgrown	1	8
prosperously	1	12
nominal	1	7
finis	1	5

[8239 rows x 2 columns]

Save Combo

Do this for safe keeping.

Students are not asked to do this, so don't worry if it's not there.

`CORPUS.to_csv(csv_combo)`

Answer Questions

1. How many raw tokens are in the combined data frame?

```
CORPUS.shape[0]
```

```
207271
```

2. How many distinct terms are there in the combined data frame (i.e. how big is the vocabulary)?

```
V.shape[0]
```

```
8239
```

3. How many more terms does the vocabulary of *Sense and Sensibility* have than that of *Persuasion*?

Method 1

```
vc_sense = CORPUS.loc[1].term_str.value_counts().shape[0]
```

```
vc_persu = CORPUS.loc[2].term_str.value_counts().shape[0]
```

```
vc_sense - vc_persu
```

```
520
```

Method 2

Students don't have to do this, but it's a good idea to put features where they belong.

In this case, we can think of the the term counts per book as features of *V*.

```
V['in_1'] = CORPUS.loc[1].term_str.value_counts()
```

```
V['in_2'] = CORPUS.loc[2].term_str.value_counts()
```

```
V.in_1.count() - V.in_2.count()
```

```
520
```

A second way to do this, which does not rely on the existences on NAs, is to convert the values to booleans.

```
V.in_1.fillna(0).astype('bool').sum() - V.in_2.fillna(0).astype('bool').sum()
```

```
520
```

4. What is the average number of tokens, rounded to an integer, per chapter in the corpus?

```
CORPUS.groupby(OHC02[:2]).term_str.count().mean().round().astype('int')
```

```
2778
```

5. What is the average number of tokens, rounded to an integer, per paragraph in the corpus?

```
CORPUS.groupby(OHC02[:3]).term_str.count().mean().round().astype('int')
```

73